

DigiT — AI Intelligence

The brain of DigiT: enterprise intelligence that perceives, reasons & acts · v1.0

AI Intelligence is not another chatbox bolted onto a module. It is a **brain** — wired into every module in DigiT, running a continuous **perceive** → **reason** → **act** loop. A chatbot is *pull*: you ask, it answers, it sleeps. A brain is a loop that never stops: it watches every module, reasons across all of them, and surfaces what matters *before* you ask. That single shift is what makes it enterprise intelligence rather than a help desk — and it fills a gap no siloed tool can, because none of them can see the whole operation at once. The positioning, in one line: **the first AI that understands your entire operation — not a corner of it — and acts on it.**




Figure 1 — The loop: every module feeds live data and events into a brain that perceives, reasons and acts — non-stop — producing real-time insight, benchmarks, predictions and resources, and acting back on the modules.

How the brain works — the three phases

Each phase maps onto an engine you are already building.

Perceive — it has senses into everything

Two channels. It **subscribes** to the cross-workspace event bus (Postgres LISTEN/NOTIFY / Supabase Realtime), so the instant anything happens — a cert lapses, a deal moves, a metric crosses a line, an asset goes quiet — the brain knows. And it has scoped **read access** to every module's tables (the same `tenant_id` + RLS boundary as everything else), so it can pull the full picture on demand. It runs both ways: reactively on events, and on a schedule (cron) for periodic sweeps.

Reason — where the scripts, commands & scores live

The brain keeps a continuously-updated **operational graph**: the entities across workspaces and the links between them (this cert → these assets → these batches → this deal → this customer → this revenue). On top of it, it runs a **library of analytical jobs** — compute the DigiT Score, run an anomaly scan, fit a forecast, trace a causal chain, run a benchmark, simulate a scenario. The LLM core decides which to run, composes them, and — critically — **ranks what it finds by business impact**, so it surfaces the three things that matter, not forty alerts.

Act — it does not just report, it does

Within guardrails you set (a confidence threshold, scoped permissions, human-in-the-loop), it surfaces the insight *and* can trigger the fix through the Workflow engine — create the task, draft the message, open the CAPA, reschedule the maintenance. Below the confidence line it asks first; above it, it acts and logs. That is

the line between a dashboard that tells you and a brain that handles it.

What you see — a feed, not a chatbox

If the brain leads, the interface cannot be an empty chat window waiting for input.

The surface is an **intelligence feed / command center**: a living, ranked stream of what the brain is currently thinking — the active watches, the causal chains it has flagged, the predictions, and the actions pending your approval — with the morning briefing at the top. You *can* still ask it things (a command bar is there), but asking is the secondary mode; its primary job is **pushing** intelligence to you. Think mission control, not messenger.

AI Intelligence — the command center (a feed, not a chatbox)





Figure 2 — The command center. The brain pushes the few things that matter, each with its causal chain and a one-tap action. The command bar is present but de-emphasised — you do not have to ask.

The hero capability — cross-workspace causal reasoning

The signature move no siloed tool can copy.

AI Intelligence connects a cause in one corner of the business to its consequence in another — all the way to the rupee. The example below touches Compliance, Resources, the CRM, Performance and the DigiT Score, and ends in a number and an action. No single-point tool can write that sentence, because none of them holds all the facts.

Cross-workspace causal reasoning — one cause, traced to the rupee



No siloed tool can write this sentence — it needs all four facts at once. That is the moat.

Figure 3 — A single lapsed certificate, traced across four workspaces to a 9-point score drop and a 14-lakh deal at risk — with the fix attached. This one paragraph is the demo, the sales deck, and the moat.

What makes it a real intelligence, not a one-trick

- **It works the night shift** — not a box you query — a brain that watches 24/7 and hands you the 3 things that matter each morning, ranked by impact. It earns trust by being right about what is important.
- **It runs the simulation before you commit** — “simulate: we take the Globex contract” → it models the load on your resources, the compliance implications, the cash flow and the team capacity, and tells you what breaks before you sign. A flight simulator for decisions.
- **It sees around corners** — it learns your operation's patterns and warns weeks ahead (“~3 weeks from a stockout that threatens the Acme contract”). The value is not the alert — it is the lead time to prevent it.
- **It can actually do the work** — within your guardrails, specialised agents act — a compliance watchdog that opens CAPAs, a follow-up agent that drafts and (on approval) sends, a maintenance scheduler. AI staff with scoped permissions and a confidence threshold.
- **You command it by intent** — “get me ready for the ISO audit next month” → it assembles the evidence, flags the gaps, creates the tasks and drafts the responses — orchestrating across workspaces from one sentence.

Why the moat deepens over time

- **It learns your specific business** — not a generic model — the longer you use it, the more it knows your suppliers, your cycles, your customers.
- **Anonymised network effect** — because DigiT is multi-tenant, it can tell a client “businesses like yours that fixed this early kept the deal 80% of the time” — a data-network effect a single-company tool can never have.
- **It always shows its work** — every insight ships with the causal chain and the underlying evidence. The chain is the explanation — which is why this design earns trust instead of demanding it. No founder acts on a black box when real money is on the line.

Build matrix — the brain is your engines, orchestrated

Not a new product — the engines you are building, pointed at the whole operation at once.

Brain capability	Powered by
Perceive — event subscription + scoped reads across modules	Cross-workspace event bus + #1 Multi-Tenant / #3 RBAC
Reason — operational graph, scoring, anomaly, forecast, causal trace, benchmark, simulate	#7 Analytics Engine
Orchestration, causal narration, intent parsing, summarisation	#5 AI Assistant
Act — create tasks, draft, trigger, autonomous agents	#6 Workflow Automation
Ask-across-everything, evidence retrieval	#4 Universal Search
Guardrails — scoped permissions, confidence threshold	#3 RBAC

In simple terms

Most products call a chatbot “AI.” DigiT's AI Intelligence is different in kind: a brain that is plugged into every module, never stops watching, connects cause to consequence across the whole operation, and acts on it — while always showing its work. It is the one feature competitors structurally cannot copy, because it depends on something only DigiT has: all the operational data, in one connected place. That is what makes it the highlight.

The one-line bar: the first AI that understands your *entire* operation — not a corner of it — perceives, reasons and acts on it non-stop, and proves every conclusion with the chain that led to it.