

DigiT — Operational Workspace

From an intake slot to the platform's intake & understanding hub · v1.0

Today the operational workspace has two doors — *upload document* and *paste text* — and whatever lands there simply sits. The reframe: this is not *where you put files*, it is **the mouth of DigiT** — the single front door where messy real-world input arrives, is understood, and is pushed out as clean data, tasks and workflows. Upload and paste are two doors; the workspace needs more doors and, above all, a **brain** behind them. The build has four layers — **ways in**, an **understanding layer**, an **action layer**, and a **working surface** — plus reusable **recipes** and the cross-platform bets only DigiT can make.




Figure 1 — Many ways in feed a single understanding & action hub that classifies, extracts and routes input out to the right workspace as clean data, tasks and workflows.

1. More ways in (beyond upload + paste)

Each is another door through which real-world input can arrive.

- **Email-to-workspace** — forward any email to a DigiT address and it is ingested and processed.
- **WhatsApp + voice notes from the field** — a plant worker forwards a voice note; it lands here transcribed.
- **Photo / scan capture** — snap an invoice, a whiteboard or a handwritten log → OCR turns it into text.
- **Audio / meeting recording** — → automatic transcript.
- **URL / web clip** — paste a link and it pulls and extracts the content.
- **Connectors** — pull straight from Drive, Gmail or Slack.
- **Bulk drop** — process 50 files at once, not one at a time.

2. The understanding layer

What happens the instant input lands — the difference between a folder and an operational workspace.

- **Auto-classify** — what is this? Invoice, contract, SOP, complaint, report, COA.
- **Structured extraction** — turn a PDF into fields: parties, dates, amounts, line items, obligations, signatures.
- **Summarize** — a TL;DR, key points and sentiment.
- **Detect intent & required action** — “this is a complaint → needs a ticket”; “this is an invoice → log a payable.”

- **Auto-route** — a compliance doc goes to Compliance; an asset doc to Resources; a cost to Performance.
 - **Confidence + source highlighting** — shows where in the document each fact came from, with a confidence score — trust matters when it is operational.
-

3. The action layer

Turn understanding into operations.

- **Auto-create** — tasks, tickets or workflow triggers from the document.
 - **Populate records** — create a vendor, asset, contact or compliance entry from what was extracted.
 - **Draft a response** — reply to the email, generate the acknowledgment.
 - **Trigger a workflow** — via the Workflow Automation engine.
 - **Human-in-the-loop** — review and approve before anything commits — essential for accuracy.
-

4. The working surface

A place to operate, not just store.

- **Unified feed** — everything ingested, with status: new → processed → actioned.
 - **Split view** — the original document alongside the extracted structured view alongside an AI chat about it.
 - **Ask across all documents** — RAG over the whole workspace corpus (“what did the March inspection flag?”).
 - **Compare / diff** — two contract versions, two SOPs, this month's report vs last.
 - **Command bar** — type or speak an instruction and it operates on the selected documents.
-

5. Recipes / playbooks

The power-user feature — save an input→action flow once, run it forever.

- **Invoice** — extract → log to payables → flag anomalies.
 - **Contract** — extract terms → set renewal reminder → file in the vault.
 - **Meeting recording** — action items → tasks + a follow-up email.
 - **Inspection report** — findings → auto-create a CAPA in Compliance.
 - **Vertical recipes** — pharma (batch records, COAs, deviation reports), renewables (site surveys, generation reports, O&M logs), logistics (PODs, e-way bills, manifests).
-

6. The creative bets only DigiT can make

- **Universal intake + auto-routing** — the connective front door for all unstructured input across the company, feeding every other workspace. No point tool does this.
 - **Vertical document intelligence** — pre-built extractors that already know what a batch record, a SECI generation report or an e-way bill looks like.
 - **Field-ops via voice / WhatsApp** — a voice note — “Plant 2 compressor down” — is transcribed, classified, turned into a maintenance task in Resources, and alerts the right person, hands-free from the floor.
 - **Document → executable workflow** — drop an SOP and it builds a runnable workflow in the Workflow engine.
 - **Operational copilot** — hand it a document plus an instruction and it executes the whole operational task end-to-end.
-

Why this is operational intelligence, not a file uploader

The operational workspace is the **ingestion + understanding + routing hub** that feeds the rest of the platform. An invoice dropped here becomes a payable, a vendor record *and* a cost data point in Performance. An inspection report becomes CAPAs in Compliance. A delivery document becomes a logistics

record. The three workspaces are where the *digested* output lives — this is where raw reality enters and gets digested.

Build matrix — almost none of this is from scratch

You are surfacing the engines you are already building through an intake lens.

Capability group	Powered by
Classify, extract, summarize, detect intent, draft, operational copilot	#5 AI Assistant
Ask-across-docs, split view, compare / diff	#4 Universal Search
Auto-route, auto-create tasks/records, recipes, doc→workflow, triggers	#6 Workflow Automation
Cross-document patterns, anomaly flags, the cost/data feed to Performance	#7 Analytics Engine
Storage, versioning, OCR pipeline, the evidence/document layer	#12 Document Management
Per-tenant isolation, role-scoped access to ingested data	#3 RBAC + #1 Multi-Tenant

Suggested build order

To make the two existing doors into a real operational hub, add these first, in order:

- 1. Understanding layer on upload + paste — classify, extract and summarize what is already coming in — this alone makes the existing doors feel intelligent.**
- 2. Action layer with human-in-the-loop — auto-create tasks / records from a document, with review before commit.**
- 3. Auto-route to the three workspaces — so digested output lands where it belongs — the connective tissue of the platform.**
- 4. Recipes / playbooks — let users save and replay the flows they run most.**
- 5. New doors (email, WhatsApp/voice, OCR) — widen intake once the brain behind it is proven.**

The sequencing point: do not add more doors first. The value is the brain behind the doors — make upload and paste intelligent, then widen intake. Field-ops voice / WhatsApp and vertical extractors are high-value but come once the core understand→route→act loop is proven.